

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [2]: train_data = pd.read_csv('train.csv')
test_data = pd.read_csv('test.csv')
```

```
In [19]: for column in train_data.columns:
    if train_data[column].dtype == 'object':
        # Fill missing values with the mode for categorical features
        train_data[column].fillna(train_data[column].mode()[0], inplace=True)
        if column in test_data.columns:
            test_data[column].fillna(test_data[column].mode()[0], inplace=True)
    else:
        # Fill missing values with the mean for numeric features
        train_data[column].fillna(train_data[column].mean(), inplace=True)
        if column in test_data.columns:
            test_data[column].fillna(test_data[column].mean(), inplace=True)

# Select features
features = ['GrLivArea', 'BedroomAbvGr', 'FullBath', 'HalfBath', 'TotRmsAbvGr']
X = train_data[features]
y = train_data['SalePrice']

# Split the training data for validation
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)


# Train the model
model = LinearRegression()
model.fit(X_train, y_train)
```

Out[19]:


```
LinearRegression()  
LinearRegression()
```

```
In [4]: y_pred = model.predict(X_val)  
mae = mean_absolute_error(y_val, y_pred)  
mse = mean_squared_error(y_val, y_pred)  
r2 = r2_score(y_val, y_pred)
```

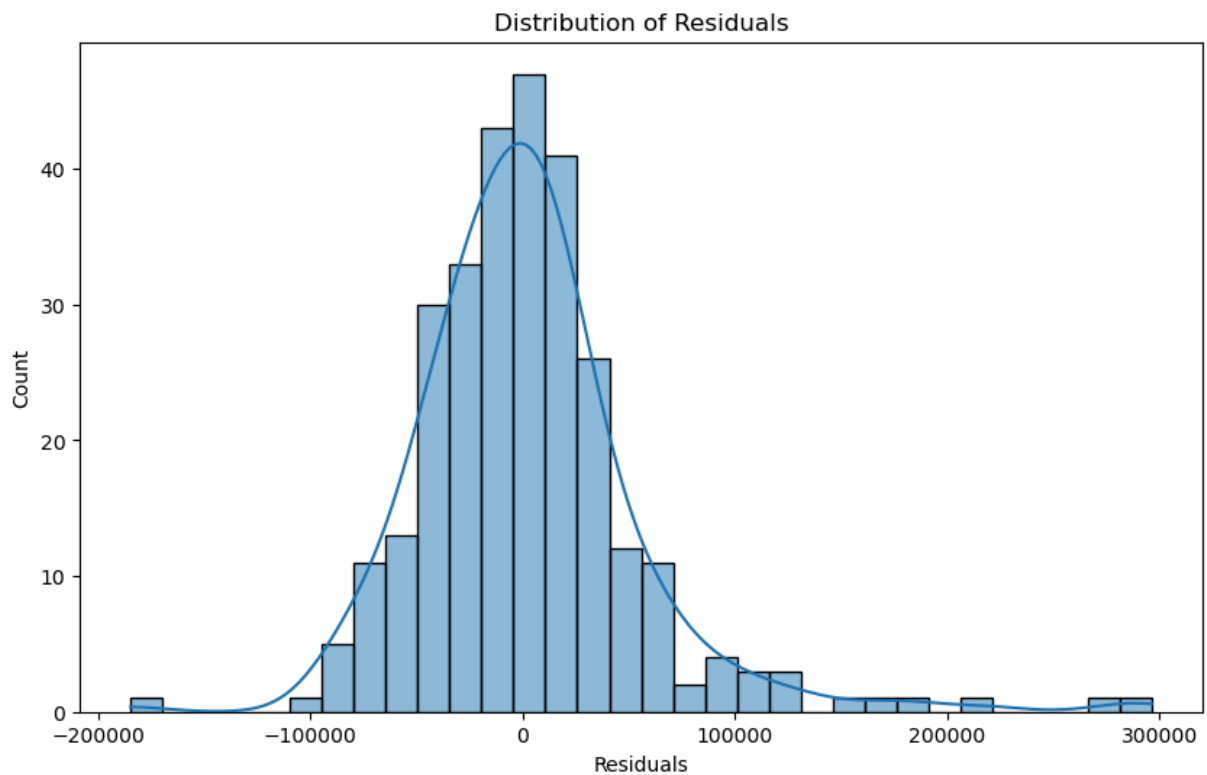
```
In [5]: plt.figure(figsize=(10, 6))  
plt.scatter(y_val, y_pred, alpha=0.5)  
plt.xlabel('Actual Sale Price')  
plt.ylabel('Predicted Sale Price')  
plt.title('Actual vs Predicted Sale Price')  
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--')  
plt.show()
```



```
In [6]: residuals = y_val - y_pred  
plt.figure(figsize=(10, 6))  
plt.scatter(y_pred, residuals, alpha=0.5)  
plt.xlabel('Predicted Sale Price')  
plt.ylabel('Residuals')  
plt.title('Residual Plot')  
plt.axhline(y=0, color='r', linestyle='--')  
plt.show()
```



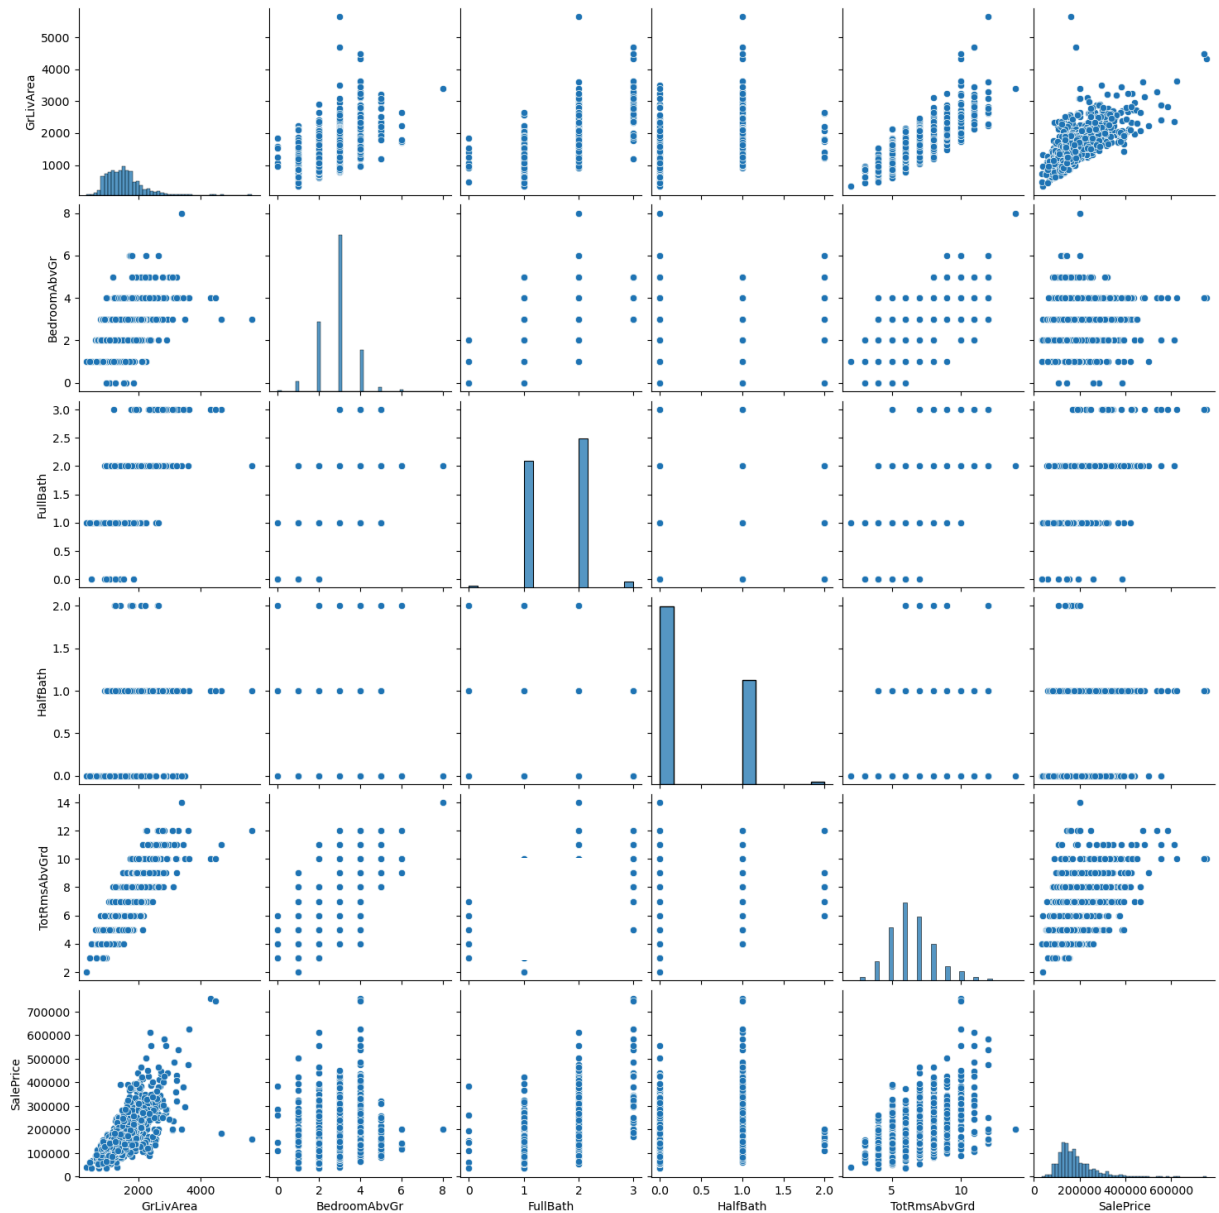
```
In [7]: plt.figure(figsize=(10, 6))
sns.histplot(residuals, kde=True)
plt.xlabel('Residuals')
plt.title('Distribution of Residuals')
plt.show()
```



```
In [8]: plt.figure(figsize=(12, 8))
sns.pairplot(train_data[features + ['SalePrice']])
```

```
plt.show()
```

<Figure size 1200x800 with 0 Axes>



```
In [21]: example = pd.DataFrame({
    'GrLivArea': [2000],
    'BedroomAbvGr': [3],
    'FullBath': [2],
    'HalfBath': [1],
    'TotRmsAbvGrd': [7]
})
example_prediction = model.predict(example)
print(f'Example Prediction: ${example_prediction[0]:,.2f}')

# Prepare the test data and make predictions
X_test = test_data[features]
test_predictions = model.predict(X_test)

# Save predictions
```

```
submission = pd.DataFrame({'Id': test_data['Id'], 'SalePrice': test_predicti  
submission.to_csv('submission.csv', index=False)
```

Example Prediction: \$240,896.28

In [ ]: